

OPERATING SYSTEM ASSIGNMENT - 4

RUCCHIKA KAPOOR

2301010240

PART-B

Q8

Distributed Deadlock Detection:

given fragments:

- $S_1 : P_1 \rightarrow P_2, P_3 \rightarrow P_4$
- $S_2 : P_2 \rightarrow P_5, P_5 \rightarrow P_6$
- $S_3 : P_6 \rightarrow P_1$

a) Global wait - for graph all

combine edges (all):

$P_1 \rightarrow P_2$

$P_2 \rightarrow P_5$

$P_5 \rightarrow P_6$

$P_6 \rightarrow P_1$ (cycle)

$P_3 \rightarrow P_4$

b) Deadlock detection

The cycle: $P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$

Hence, deadlock exists. Process
processors in deadlock:

P_1, P_2, P_5, P_6

c) Distributed algorithm to detect it

A suitable algorithm:

Chandy - Misra - Haas Probe - Based Deadlock

Detection sends Algorithm

This algorithm sends PROBE messages across sites to detect cycles in distributed wait for graph.

Q4

Distributed file system performance:

given:

$$\text{Local access} = 5 \text{ ms}$$

$$\text{Remote access} = 25 \text{ ms}$$

$$\text{Probability} = 0.3 \rightarrow \text{local} = 0.4$$

a) Expected access time:

$$E = (0.4 \times 5) + (0.3 \times 25)$$

$$= 3.5 + 7.5 = 11 \text{ ms}$$

b) Caching strategy:

Client-side caching with LRU (Least Recently Used)

JUSTIFYING:

Users usually access the same files repeatedly. Caching them locally reduces remote accesses → significantly lowering avg. access time and network load.

Q5

Check Pointing Strategy:

given:

$$\text{full checkpoint} = 200 \text{ ms}$$

$$\text{incremental} = 50 \text{ ms}$$

$$\text{RPO} = 1 \text{ sec}$$

$$\text{period} = 10 \text{ sec}$$

a) Optimal mix:

- Take 1 full checkpoint at the start
- then take incremental checkpoints every 1 second → total 9 incremental checkpoints

Total overhead :

full : 200 ms

incremental : $9 \times 50 = 450$ ms

total : 650 ms over 10 seconds

b) Reasoning :

$\rightarrow RPO = 1s \rightarrow$ system must checkpoints at least once every 1 second

- \rightarrow full checkpoint is expensive; incremental is cheaper
- \rightarrow taking only incremental checkpoints avoids high overhead while staying without RPO.
- \rightarrow one full checkpoint ensures a clean baseline for recovery and consistency.

Q9 Case Study: Global E-commerce Platform

- a) distributed scheduling flash sales:
- \rightarrow sudden traffic spikes across continents
- \rightarrow hotspot servers overloaded while others idle.
- \rightarrow maintaining fairness across regions.
- \rightarrow handling network latency & replication delays.
- \rightarrow Ensuring fast task migration b/w servers

Suitable load-balancing algorithm :

consistent hashing with dynamic load redistribution.

BEST CHOICE:

Work stealing algorithm Idle servers "steal" tasks from overloaded servers \rightarrow perfect for unpredictable flash sale workloads

b) Fault tolerance strategy (RTO and RPO focused)

use geo-replicated active-active clusters with:

- synchronous replication inside a region (low RPO)
- asynchronous cross-region replication (low latency)
- automatic failover using health checks + DNS based global load balancing.

ENSURES:

→ low RPO:

Data locs minimised since updated are replicated frequently.

→ low RTO:

Traffic can switch to a healthy region within seconds.

EXAMPLES:

- multi-region distributed file system.
- checkpoint + log-based recovery in each region.
- stateless microservices + replicated databases.

Getting
2 ||| 25